

Quick Intro to Version Control Systems and Project Hosting Services

Hans Petter Langtangen^{1,2}

¹Center for Biomedical Computing, Simula Research Laboratory,

²Department of Informatics, University of Oslo

Aug 11, 2013

1 Introduction

Version control systems allow you to record the history of files and share files among several computers and collaborators in a professional way. Especially when working with software development or technical reports it is essential to have changes documented and to ensure that every computer and person involved in the project have the latest updates of the files. Greg Wilson's excellent [Script for Introduction to Version Control](#) provides a more detailed motivation why you will benefit greatly from using version control systems.

Projects that you want to share among several computers or project workers are today most conveniently stored at some web site "in the cloud" and updated through communication with that site. I strongly recommend you to use such sites for all serious programming and scientific writing work – and all other important files.

1.1 Why not Dropbox or Google Drive?

The simplest services for hosting project files are [Dropbox](#) and [Google Drive](#). It is very easy to get started with these systems, and they allow you to share files among laptops and mobile units with as many users as you want. The systems offer a kind of version control in that the files are stored frequently (several times per minute), and you can go back to previous versions for the last 30 days. However, it is challenging to find the right version from the past when there are so many of them and when different versions are not annotated with sensible comments.

More seriously, when several people may edit files simultaneously, it can be difficult to detect who did what when, roll back to previous versions, and to manually merge the edits when these are incompatible. Then one needs more sophisticated tools than Dropbox or Google Drive: project hosting services with true version control systems. The following text aims at providing you with the

minimum information to started with such systems. Numerous other tutorials contain more comprehensive material and in-depth explanations of the concepts and tools.

Essence.

The essence with project hosting services is that you have the files associated with a project in the cloud. Many people may share these files. Every time you want to work on the project you explicitly update your version of the files, edit the files as you like, and synchronize the files with the "master version" at the site where the project is hosted. If you at some point need to go back to a version of the files at some particular point in the past, this is an easy operation. You can also use tools to see what various people have done with the files in the various versions.

1.2 Project hosting services

Four popular project hosting services are

- [Bitbucket](https://bitbucket.org) at bitbucket.org
- [GitHub](https://github.com) at github.com
- [Googlecode](https://code.google.com) at code.google.com
- [Launchpad](https://launchpad.net) at launchpad.net
- [Sourceforge](https://sourceforge.net) at sourceforge.net

All these services are very similar. Below we describe how you get started with Bitbucket, GitHub, and Googlecode. Launchpad works very similarly to the latter three. All the project hosting services have excellent introductions available at their web sites, but the recipes below are much shorter and aim at getting you started as quickly as possible by concentrating on the most important need-to-know steps.

1.3 Repositories and local copies

The mentioned services host all your files in a specific project in what is known as a *repository*, or *repo* for short. When a copy of the files are wanted on a certain computer, one clones the repository on that computer. This creates a local copy of the files. Now files can be edited, new ones can be added, and files can be deleted. These changes are then brought back to the repository. If users at different computers synchronize their files frequently with the repository, most modern version control systems will be able to merge changes in files that have been edited simultaneously on different computers. This is perhaps one of the most useful features of project hosting services. However, the merge functionality clearly works best for pure text files and less well for binary files, such as PDF files, MS Word or Excel documents, and OpenOffice documents.

1.4 What to choose?

It might seem challenging to pick the project hosting service and the version control system that are right for you. Below is some very personal advice.

- I would recommend Git as the default choice of version control system.
- If you need private repos and you are not a student, choose Bitbucket unless you are willing to pay for private repos at GitHub.
- If you want to make links to your programs in documentation, stay away from Bitbucket as this site embeds the version of the file into its URL (but you can get around this problem with a trick described in Section 2.3).
- If you want to easily drop HTML files for documentation around in your repo, Googlecode will display them as web pages right away, while the procedure on GitHub is technically more complicated (Bitbucket does not yet support web pages as part of a project). Unfortunately, Googlecode does not allow MathJax for rendering L^AT_EX mathematics in HTML files.
- If a few collaborate on a project and everyone wants to be notified in email about changes in project files, Bitbucket and Googlecode have better support for this feature than GitHub.

For the examples below, assume that you have some directory tree `my-project` with files that you want to host at Bitbucket, GitHub, or Googlecode and bring under version control. The official name of the project is "My Project".

1.5 Installing Git

The installation of Git on various systems is described on the [Git website](#) under the *Download* section. Git involves compiled code so it is most convenient to download a precompiled binary version of the software on Windows, Mac and other Linux computers. On Ubuntu or any Debian-based system the relevant installation command is

```
Terminal> sudo apt-get install git gitk git-doc
```

This tutorial explains Git interaction through command-line applications in a terminal window. There are numerous [graphical user interfaces to Git](#). Three examples are [git-cola](#), [TortoiseGit](#), and [SourceTree](#).

Make a file `.gitconfig` in your home directory with information on your full name, email address, your favorite text editor, and the name of an "excludes file" which defines the file types that Git should omit when bringing new directories under version control. Here is a simplified version of the author's `.gitconfig` file:

```
[user]
name = Hans Petter Langtangen
email = hpl@simula.no
editor = emacs

[core]
excludesfile = ~/.gitignore
```

The "excludes file" is called `.gitignore` and must list, using the Unix Shell Wildcard notation, the type of files that you do not need to have under version control, because they represent garbage or temporary information, or they can easily be regenerated from some other source files. A suggested `.gitignore` file looks like

```
# compiled files:
*.o
*.so
*.a
# temporary files:
*.bak
*.swp
*~
.*~
*.old
tmp*
temp*
# tex files:
*.log
*.dvi
*.aux
*.blg
*.idx
# eclipse files:
*.cproject
*.project
# misc:
.DS_Store
```

Be critical to what kind of files you really need a full history of. For example, you do not want to populate the repository with big graphics files of the type that can easily be regenerated by some program.

2 Bitbucket

To start using Bitbucket, go to bitbucket.org and create an account. The communication channel with Bitbucket repositories is either through SSH or HTTPS. To use SSH, you must upload your SSH key, typically the contents of the file `id_rsa.pub` or `id_dsa.pub` in the `.ssh` subdirectory of your home directory. Go to the page for your account, choose *SSH keys*, and upload one of these files. The essence of the SSH keys is that they allow you to communicate with Bitbucket without using a password, which is very convenient.

There are links to extensive help pages if you do not have such keys or if you are unfamiliar with SSH. Follow these steps on Mac or Linux machines to generate keys: 1) check that you have `ssh` installed; 2) create a `.ssh` directory in your home directory; and 3) run `ssh-keygen` in the `.ssh` directory (you are prompted for a passphrase - just write something). On Windows one applies the PuTTY and the TortoiseHG programs to generate and register keys, see the help pages on Bitbucket. Once the keys are generated, you can continue using them on any computer.

2.1 Creating a new project

Click at *Repositories* and at *Create repository*. You can now

- fill in the name of the project, here `my-project`,
- decide whether the project is private or public (the number of private repos is unlimited for yourself, but you have to pay to invite more than five users in total to share your private repos unless you are certified to be a student or work at an academic institution),
- choose between the Git or Mercurial version control system (here we assume you run Git),
- click whether you want issue tracking for reporting errors, suggesting improvements, etc.,
- click whether you want a wiki page associated with the project,
- fill in a brief description,
- click on *Create repository*.

While doing this you may also want to have the [Bitbucket 101 guide](#) available.

It is now time to *clone* (copy) the project to your laptop. Go to the project page (you are automatically brought to the project page when creating a new project). Find the *Clone* button, click on it, choose *SSH*, copy the *clone* line and run this command in a terminal:

```
Terminal> git clone git@bitbucket.org:user/my-project.git
```

You must replace `user` with your own username at Bitbucket and `my-project` by the real project name.

The first time you do the clone command you may be prompted by cryptic Unix output ending with "Are you sure you want to continue connecting (yes/no)?" . Just answer yes.

The next step is to collect files and directories that should make up the project and put them in the `my-project` directory. Standing in the `my-project` directory, the following Git command is used to add all files in the current directory *tree*, except those having file types listed in `.gitignore`:

```
Terminal> git add .
```

Thereafter, the changes to the repository (here adding of files) must *committed* (registered):

```
Terminal> git commit -am 'Initial import of files.'
```

The text following the `-am` option is a required description of the changes that have taken place. This description does not matter much for this initial import of files into the repository, but is of importance for future commit commands so that you can easily track the history of your project files and inform others (and yourself) what you have done with the files.

The final step is to push the local changes to the master repo at Bitbucket:

```
Terminal> git push -u origin master
```

You must be connected to the Internet for the `push` command to work since it sends file information to the `bitbucket.org` site.

Further work with the files must always follow the pull, edit, commit, and push steps explained in Section 5 for Git.

Collaborating. There is a button *Send invitation* on the project home page where you can invite other Bitbucket users to have push (write) access to your repo. Many prefer to be notified in email when changes are pushed to the repo: click the settings wheel to the right, choose *Services*, then *Email*, and fill in the email addresses that are to receive notifications on updates. Under *Access management* you can fill in the Bitbucket names of users who are have pull (read) or push (write) access to the repo.

2.2 User web pages

Bitbucket can host web pages associated with a Bitbucket account, but does not yet offer web pages as part of a project. Say your account/username is `user`. Make a new repository on Bitbucket called `user.bitbucket.org`. Clone it, fill it with a file, and push back, e.g.,

```
Terminal> git clone git@bitbucket.org:user/user.bitbucket.org.git
Terminal> cd user.bitbucket.org
Terminal> echo "Welcome to my web pages!" > test.html
Terminal> git add test.html
Terminal> git commit -am 'First web page'
Terminal> git push origin master
```

You can now load the URL `http://user.bitbucket.org/test.html` into a web browser and view the page. By creating various subdirectories you can host web pages for a series of projects in this repo.

2.3 Linking to Bitbucket files

Unfortunately, the default URL shown in your browser to a file in a Bitbucket repo contains information about the version of the file. For example, a file `v1.c` under Git might appear with the URL

```
https://bitbucket.org/user/proj/src/5fb228107044/dir/v1.c?at=master
```

The string `5fb228107044`, called commit hash, is connected to the version of this file and will change when the file is updated. Such varying URLs are useless when you to link to the file in documents. To make a stable link to the latest version of a file (in a public repo), replace the commit hash by `master` and remove the final `?at=master`:

```
https://bitbucket.org/user/proj/src/master/dir/v1.c
```

Other endings of the URL with commit hash, e.g., `?at=default`, requires the hash to be replaced by `default`.

3 Googlecode

To use Googlecode you need a general account on Google, which allows you to use Gmail, Google Docs, and other Google services. Googlecode offers free file hosting, but all the files are public and open to the world.

3.1 Creating a new project

Go to <http://code.google.com/hosting> and click on *Creating a new project*. Fill out *all* the fields. For now the project name is **my-project**. You can choose between Git, Mercurial, or Subversion as version control system for your project, and this choice is difficult to change later. If you choose Git, you must create a file `.netrc` in your home directory containing the line

```
machine code.google.com login uname password pw
```

Here, `uname` is your username for the Google account and `pw` is the Googlecode password which is generated for you on <https://code.google.com/hosting/settings> when you are logged in. The `.netrc` file avoids typing or pasting in your long and complicated password every time you push changes to the repository on `code.google.com`.

The next step is to clone the empty repository on your local machine so that you can add files:

```
Terminal> git clone https://code.google.com/p/my-project/
```

Now you can go to the **my-project** directory and add files. Perform

```
Terminal> git add .
Terminal> git commit -am 'First import of files.'
Terminal> git push origin master
```

Click on *Source* and *Browse* on the project's web page, and observe that the added files are visible on the project page.

A very strong and useful feature with Googlecode, in my opinion, is that one can place documentation of the project anywhere in the repository, and HTML and PDF will be correctly shown in web browsers. (All other project hosting sites demands either wiki pages or special web areas for locating documentation.) The URL to your files is

```
https://my-project.googlecode.com/git/
```

For other version control systems just replace `git` by `hg` or `svn`. For example, if we have HTML documentation of our project in the directory `doc/html`, we can point users to

```
https://my-project.googlecode.com/git/doc/html/index.html
```

The HTML will be rendered correctly as opposed to when you load the similar file into the web browser from the repository,

```
http://code.google.com/p/my-project/source/browse/doc/html/index.html
```

Now you can only see the HTML source code of this file, as is usual on other project hosting sites. We remark that MathJax mathematics within the HTML code is *not* rendered correctly. At the time of this writing, GitHub is the only service that offers full MathJax support when you need web pages with mathematics.

You can click on *Project Home* and then on *Administer* to edit the main page of the project. This is a wiki, using Google's wiki syntax, but it is quite easy to add links to your documentation, e.g.,

Browse the
[<https://my-project.googlecode.com/git/doc/API/html/index.html>
API documentation].

It is easy to allow others to push their changes to the repository: click on *Sharing* and then on *Administer*. The Google account names of people you allow push (write) access can be listed under each other in the *Project committers* field.

3.2 Wiki pages

Wiki pages can intuitively be made directly in the browser. However, it is often more convenient to have them locally on your computer. Click on *Source* and choose *wiki* on the *Repository* pull down menu. The proper clone command to get a copy of the wiki repository then appears.

Googlecode applies their own Google wiki format. My preference is to write documentation in the neutral [Doconce](#) format and transform the document to L^AT_EX, HTML, Sphinx, and Google wiki. The wiki can then be copied from the project directories to the wiki directory and then pushed to the repository. This ensures that there is only one source of the documentation (despite the need for many formats) and that the wiki pages are frequently updated.

4 GitHub

Go to github.com and create an account. Then go to your account settings (icon in the upper left corner of the page), choose *SSH Keys*, and provide your SSH key unless you have already registered this key with another GitHub account (see Appendix A). Often, it is just a matter of pasting the contents of `id_rsa.pub` or `id_dsa.pub` files, located in the `.ssh` subdirectory of your home directory, into the *Key* box in the web page. Make sure to just cut and paste the text from, e.g., `id_rsa.pub` without any extra whitespaces or other text. How to generate these files is described in the link *generating SSH keys* above the SSH Keys box (or see the introduction to Bitbucket above).

If the account is a project account and not a personal account, I do not recommend to provide an SSH key although it can be done (see Appendix A). It is easier to log in and add collaborators using their personal GitHub usernames.

4.1 Creating a new project

Click on *New repository* on the main page and fill out a project name, here *My Project*, click the check button *Initialize this repository with a README*, and click on *Create repository*. Unless you pay, all repos are public, but students and teachers can request [free](#), [private](#) repos.

The next step is to clone the project on your personal computer. Click on the *SSH* button to see the address of the project, and paste this address into a terminal window, after `git clone`:

```
Terminal> git clone git://github.com:user/My-Project.git
```

Make sure you substitute `user` by your own username on GitHub.

The result of the `git clone` command is a new directory `My-Project`. It contains the file `.git`, which shows that it is a Git repository. It also contains a default `README.md` file with the project name and description. The extension `.md` signifies a file written in the [Markdown](#) format. You may use the [reStructuredText](#) format as an alternative (`README.rst`), or simply write a plain text file (`README`), but the `git mv` command must be used to change the filename.

You can now add files and directories into the `My-Project` directory. When your initial file collection has the desired form, you must run

```
Terminal> git add .
Terminal> git commit -am 'First set of files.'
Terminal> git push -u origin master
```

The daily file operations are explained in [Section 5](#).

Collaborating. To give others permissions to push their edits of files to the repository, you click on the *Settings* link in the right sidebar, then click on *Collaborators* on the left, and fill in the name of a collaborator (her or his username on GitHub). Many find it convenient to be notified in email when others have pushed a new version of the files to the repo. Click on *Service Hooks* (in the project's *Settings* menu), choose *Email*, fill in at most two whitespace-separated email addresses, and mark the *Active* box. More addresses must be dealt with through a mailing list and filling in the name of that list.

Anyone who participates in a project (has write access) or watches a project (having clicked the *watch* button) can monitor the development of the activity on their GitHub main page. Go to *Account Settings* and choose *Notification Center*. There you see two sections, *Participating* and *Watching*, for those participating in the project (granted write access) and those watching the project (having clicked the *watch* button), respectively.

4.2 Wiki pages

With every GitHub project there is an option to create wiki pages. Click on the *Wiki* button in the right sidebar on the main page of the project. Click on *New Page* to create a new page. The wiki pages can be written in different markup languages. Markdown is the default choice, but you can alternatively use MediaWiki (the LaTeX-supported markup used for Wikipedia) and reStructuredText (unfortunately not the Sphinx extension with nice $\text{L}^{\text{A}}\text{T}_{\text{E}}\text{X}$ mathematics).

The wiki pages can be written and maintained through the web browser interface, but it is usually more convenient to clone them on your computer as this makes it easy to add figures and other documents you may want to link to. It also makes it straightforward to edit the wiki text in your favorite text editor. The wiki pages are stored in a separate repo and can be cloned by

```
Terminal> git clone git://github.com/user/My-Project.wiki.git
```

This command makes a local copy of the pages in the directory `My-Project.wiki`, which you may prefer to have at the same level as the project directory itself in your directory tree.

Each wiki page has its own file, where the extension reflects the markup language used, e.g., `.md` for Markdown, `.rest` for reStructuredText, `.mediawiki` for MediaWiki, and `.creole` for Creole wiki. The wiki files are handled as other files in a GitHub project, i.e., you need to pull before editing and then perform commit and push. After the push you can reload the page in the web browser to monitor the effect.

You may consider having the original text in `doconce` format and generate the wiki in the reStructuredText or MediaWiki format.

Do changes, commit the usual way, and push by

```
Terminal> git push git@github.com:user/My-project.wiki.git
```

The address can be stored as `url` in `.git/config` in the root directory of the wiki project so that just a `git push` works.

4.3 Project web pages

HTML pages stored in your repo cannot be linked to and properly rendered as web pages. Links to HTML files at GitHub will be either rendered as nicely typeset HTML code or as plain text.

Quite often you want to create a set of HTML pages to document your project. This can be done as part of your repo, but you must apply a special technique consisting in creating a separate *gh-pages* branch in Git and making the HTML files in this particular branch. The recipe is described in detail below.

1. Go to the project page on github.com and click *Settings*.
2. Click on *Automatic Page Generator* under the *GitHub Pages*.
3. Proceed clicking *Continue to Layouts*, choose a design of the `index.html` page that GitHub will create for you, and click *Publish*.
4. Go to the root directory of the project, `My-Project` and run `git fetch origin`.
5. Run `git checkout gh-pages`.

You have now a *new branch* called `gh-pages` of your project containing an `index.html` file and directories for JavaScript programs and CSS style sheets in the root directory. The `gh-pages` branch will also all files *not* contained in the *master branch*, typically redundant files you have generated and which should not be stored in the version control system (remove these manually with `git rm`). You can populate the root directory and subdirectories of your `gh-pages` branch with HTML and other files as you like. The key issue is that the people out there will only see the web pages that correspond to your HTML files in the `gh-pages` branch!

The `index.html` page is invoked by the web address

`http://user.github.io/My-Project/index.html`

where `user` is the GitHub username and `My-Project` is the project name.

The web pages and project files are now in two different branches. To see the branches, type `git branch` and the one you are in will be marked with `*` in the output. Switching to the master branch is done by `git checkout master`. Similarly, `git checkout gh-pages` switches to the gh-pages branch.

My personal preference is to have the master and gh-pages synchronized, at least in projects where I want to link to various source code files or other files from the web documentation. Sometimes I also update files in the gh-pages branch without remembering to switch to the master branch. To this end, one needs to *merge* the branches, i.e., automatically edit files in the current branch such that they are up-to-date and identical to files in another branch.

To merge the current branch with some branch named `otherbranch`, run

```
Terminal> git merge otherbranch
```

Git applies smart algorithms that very often manage to merge the files without human interaction. However, occasionally these algorithms are not able to resolve conflicts between two files. A message about the failure of the merge is seen in the terminal window, and the corresponding files have markers in them showing which sections that needs manual editing to resolve the conflicts. Run `git diff` to show the problems (you can tailor this command to your needs as explained in Section 5.4). After a manual edit, do `git commit -a`. More details on merging appears in Section 5.2.

If you want to keep the master branch and the gh-pages branch synchronized, start with merging the gh-pages branch with the master branch and push the complete file collection to the gh-pages branch. Then switch to the master branch and merge with gh-pages so you get the autogenerated `index.html` file and associated files and directories for web pages in the root directory of the master branch as well:

```
Terminal> git merge master
Terminal> touch .nojekyll
Terminal> git push origin gh-pages
Terminal> git checkout master
Terminal> git merge gh-pages
```

You *must* add an empty file `.nojekyll` in the top directory of the project pages if you want to use Sphinx-generated HTML pages (or other pages using javascripts, style sheets, and images in subdirectories whose names start with an underscore).

You can now add the documentation to the project files and maintain them in the master branch. Before publishing documents online, make sure to update the gh-pages branch by

```
Terminal> git commit -am 'Ensure commit of master branch'
Terminal> git push origin master
Terminal> git checkout gh-pages
Terminal> git merge master
Terminal> git push origin gh-pages
Terminal> git checkout master
```

Personally, I like to move the generated `index.html` file and all associated scripts, stylesheets, and images from the root directory to some more isolated place, say `doc/web`:

```
Terminal> git mv index.html params.json stylesheets/ images/ \
           javascripts/ doc/web/
```

The URL of the `index.html` file is

<http://user.github.io/My-Project/doc/web/index.html>

Linking to source code files or other files in the project is easy: just find the file in GitHub's web interface, choose which version of the file you want to link to (nicely HTML formatted version or the raw file), right-click on the link, choose *Copy Link*, and paste the link into the document you want. You can test that the link works by the Unix command `curl -O <link>`. Note that the link to a file is different from the source file's intuitive path in the repository. Typically, a source file `dir/f.py` in project `prj` is reached through

<https://github.com/user/prj/blob/master/dir/f.py?raw=true>

Sometimes you want to link to another HTML file, PDF file, movie file, or a file that is to be interpreted as a web resource by the browser. Do not use the path to the file in the repo as explained above as it will just bring the reader to the repo page. Instead, make sure the file is in the `gh-pages` branch and use a local link, like `../doc.pdf`, or the complete `gh-pages` URL to the file, say

<http://user.github.com/My-Project/doc/misc/doc.pdf>

Tip.

The ordinary GitHub URL of image files can be used in web pages to insert images from your repo, provided the image files are in the *raw* format - click the *Raw* button when viewing a file at github.com and use the corresponding URL in the `img` tag in the HTML code.

4.4 User web pages

GitHub also allows you to create user pages and organization pages not tied to any specific project. Your personal site has address <http://user.github.com>. Go to your home page on github.com and click *New repository*, and give it the project name `user.github.com`. Then follow the instructions that come up:

```
Terminal> mkdir user.github.com
Terminal> cd user.github.com
Terminal> git init
Terminal> # make an index.html file with some test text
Terminal> git add index.html
Terminal> git commit -m 'First commit'
Terminal> git remote add origin \
git@github.com:user/user.github.com.git
Terminal> git push -u origin master
```

Go to <http://user.github.com> and see how the `index.html` is rendered. You can now add various contents as in any ordinary Git repository. If you want to use Sphinx generated HTML pages, recall to add an empty file `.nojekyll`.

5 Using Git

Most Mac and Linux users prefer to work with Git via commands in a terminal window. Windows users may prefer a graphical user interface (GUI), and there are many [options](#) in this respect. There are also GUIs for Mac users. Here we concentrate on the efficient command-line interface to Git.

5.1 Basic Git commands

Cloning. You get started with your project on a new machine, or another user can get started with the project, by running

```
Terminal> git clone git@github.com:user/My-Project.git
Terminal> cd My-Project
ls
```

Recall to replace `user` by your real username and `My-Project` by the actual project name.

The pull-change-push cycle. The typical work flow with the "My Project" project starts with updating the local repository by going to the `My-Project` directory and writing

```
Terminal> git pull origin master
```

You may want to do `git fetch` and `git merge` instead of `git pull` as explained in Section 5.4, especially if you work with branches.

You can now edit files, add new files, delete files, or rename files. Typical commands are

```
Terminal> git add file1 file2.* dir1 dir2
Terminal> git rm filename
Terminal> git rm -r dirname
Terminal> git mv oldname newname
Terminal> git mv oldname ../newdir
```

When your chunk of work is ready, it is time to commit your changes (note the `-am` option):

```
Terminal> git commit -am 'Description of changes.'
```

If typos or errors enter the message, the `git commit --amend` command can be used to reformulate the message. Running `git diff` prior to `git commit` makes it easier to formulate descriptive commit messages since this command gives a listing of all the changes you have made to the files since the last commit or pull command.

You may perform many commits, to keep track of small changes, before you push your changes to the global repository:

```
Terminal> git push origin master
```

It is recommended to push and pull frequently if the work takes place in several clones of the repo (i.e., there are many users or you work with the repo on different computers). Infrequent push and pull easily leads to merge problems (see Section 5.2). Also remember that others cannot get your changes before they are pushed.

To see the status of files, run

```
Terminal> git status -s
```

Files are marked with different symbols, e.g., A for added, M for modified, R for renamed, and ?? for not being registered in the repo.

Viewing the history of files. A nice graphical tool allows you to view all changes, or just the latest ones:

```
Terminal> gitk --all
Terminal> gitk --since="2 weeks ago"
```

You can also view changes to all files, some selected ones, or a subdirectory:

```
Terminal> git log -p                # all changes to all files
Terminal> git log -p filename      # changes to a specific file
Terminal> git log --stat --summary  # compact summary
Terminal> git log --stat --summary subdir
```

Adding `--follow` will print the history of file versions before the file got its present name.

To show the author who is responsible for the last modification of each line in the file, use `git blame`:

```
Terminal> git blame filename
Terminal> git blame --since="1 week" filename
```

A useful command to see the history of who did what, where individual edits of words are highlighted (`--word-diff`), is

```
git log -p --stat --word-diff filename
```

Removed words appear in brackets and added words in curly braces.

Looking for when a particular piece of text entered or left the file, say the text `def myfunc`, one can run

```
Terminal> git log -p --word-diff --stat -S'def myfunc' filename
```

This is useful to track down particular changes in the files to see when they occurred and who introduced them. One can also search for regular expressions instead of exact text: just replace `-S` by `-G`.

Retrieving old files. Occasionally you need to go back to an earlier version of a file, say its name is `f.py`. Start with viewing the history:

```
Terminal> git log f.py
```

Find a commit candidate from the list that you will compare the present version to, copy the commit hash (string like `c7673487...`), and run

```
Terminal> git diff c7673487763ec2bb374758fb8e7efefa12f16dea f.py
```

where the long string is the relevant commit hash. You can now view the differences between the most recent version and the one in the commit you picked (see Section 5.4 for how to configure the tools used by the `git diff` command). If you want to restore the old file, write

```
Terminal> git checkout c7673487763ec2bb374758fb8e7efefa12f16dea f.py
```

To go back to another version (the most recent one, for instance), find the commit hash with `git log f.py`, and do `git checkout <commit hash> f.py`.

If `f.py` changed name from `e.py` at some point and you want `e.py` back, run `git log --follow f.py` to find the commit when `e.py` existed, and do a `git checkout <commit hash> e.py`.

Often you just need to *view* the old file, not replace the current one by the old one, and then `git show` is handy. Unfortunately, it requires the full path from the root git directory:

```
Terminal> git show \
c7673487763ec2bb374758fb8e7efefa12f16dea:dir1/dir2/f.py
```

5.2 Merging files with Git

The `git pull` command fetches new files from the repository and tries to perform an automatic merge if there are conflicts between the local files and the files in the repository. Alternatively, you may run `git fetch` and `git merge` to do the same thing as described in Section 5.4. We shall now address what to do if the merge goes wrong, which occasionally happens.

Git will write a message in the terminal window if the merge is unsuccessful for one or more files. These files will have to be edited manually. Merge markers of the type `>>>>`, `====`, and `<<<<` have been inserted by Git to mark sections of a file where the version in the repository differ from the local version. You must decide which lines that are to appear in the final, merged version. When done, perform `git commit` and the conflicts are resolved.

Graphical merge tools may ease the process of merging text files. You can run `git mergetool --tool=meld` to open the merge tool `meld` for every file that needs to be merged (or specify the name of a particular file). Other popular merge tools supported by Git are `araxis`, `bc3`, `diffuse`, `ecmerge`, `emerge`, `gvimdiff`, `kdiff3`, `opendiff`, `p4merge`, `tkdiff`, `tortoisemerge`, `vimdiff`, and `xxdiff`.

Below is a Unix shell script illustrating how to make a global repository in Git, and how two users clone this repository and perform edits in parallel. There is one file `myfile` in the repository.

```
#!/bin/sh
# Demo script for exemplifying git and merge

rm -rf tmp1 tmp2 tmp_repo    # Clean up previous runs

mkdir tmp_repo    # Global repository for testing
cd tmp_repo
git --bare init --shared
cd ..

# Make a repo that can be pushed to tmp_repo
mkdir _tmp
cd _tmp
cat > myfile <<EOF
This is a little
test file for
exemplifying merge
of files in different
```

```

git directories.
EOF
git init
git add .    # Add all files not mentioned in ~/.gitignore
git commit -am 'first commit'
git push ../tmp_repo master
cd ..
rm -rf _tmp

# Make a new hg repositories tmp1 and tmp2 (two users)
git clone tmp_repo tmp1
git clone tmp_repo tmp2
# Change myfile in the directory tmp1
cd tmp1
# Edit myfile: insert a new second line
perl -pi -e 's/a little\n/a little\ntmp1-add1\n/g' myfile
# Register change in local repository
git commit -am 'Inserted a new second line in myfile.'
# Look at changes in this clone
git log -p
# or a more compact summary
git log --stat --summary
# or graphically
#gitk
# Register change in global repository tmp_repo
git push origin master
cd ..

# Change myfile in the directory tmp2 "in parallel"
cd tmp2
# Edit myfile: add a line at the end
cat >> myfile <<EOF
tmp2-add1
EOF
# Register change locally
git commit -am 'Added a new line at the end'
# Register change globally
git push origin master
# Error message: global repository has changed,
# we need to pull those changes to local repository first
# and see if all files are compatible before we can update
# our own changes to the global repository.
# git writes
#To /home/hpl/vc/scripting/manu/py/bitgit/src-bitgit/tmp_repo
# ! [rejected]        master -> master (non-fast-forward)
#error: failed to push some refs to ...

git pull origin master
# git writes:
#Auto-merging myfile
#Merge made by recursive.
# myfile |      1 +
# 1 files changed, 1 insertions(+), 0 deletions(-)
cat myfile # successful merge!
git commit -am merge
git push origin master
cd ..

# Perform new changes in parallel in tmp1 and tmp2,
# this time causing hg merge to fail

```



```

# Change myfile in the directory tmp1
cd tmp1
# Do it all right by pulling and updating first
git pull origin master
# Edit myfile: insert "just" in first line.
perl -pi -e 's/a little/tmp1-add2 a little/g' myfile
# Register change in local repository
git commit -am 'Inserted "just" in first line.'
# Register change in global repository tmp_repo
git push origin master
cd ..

# Change myfile in the directory tmp2 "in parallel"
cd tmp2
# Edit myfile: replace little by modest
perl -pi -e 's/a little/a tmp2-replace1\ntmp2-add2\n/g' myfile
# Register change locally
git commit -am 'Replaced "little" by "modest"'
# Register change globally
git push origin master
# Not possible: need to pull changes in the global repository
git pull origin master
# git writes
#CONFLICT (content): Merge conflict in myfile
#Automatic merge failed; fix conflicts and then commit the result.
# we have to do a manual merge
cat myfile
echo 'Now you must edit myfile manually'

```

You may run this file `git_merge.sh` named by `sh -x git_merge.sh`. At the end, the versions of `myfile` in the repository and the `tmp2` directory are in conflict. Git tried to merge the two versions, but failed. Merge markers are left in `tmp2/myfile`:

```

<<<<<< HEAD
This is a tmp2-replace1
tmp2-add2

=====
This is tmp1-add2 a little
>>>>>> ad9b9f631c4cc586ea951390d9415ac83bcc9c01
tmp1-add1
test file for
exemplifying merge
of files in different
git directories.
tmp2-add1

```

Launch a text editor and edit the file, or use `git mergetool`, so that the file becomes correct. Then run `git commit -am merge` to finalize the merge.

5.3 Git working style with branching and stashing

Branching and stashing are nice features of Git that allow you to try out new things without affecting the stable version of your files. Usually, you extend and modify files quite often and perform a `git commit` every time you want to record the changes in your local repository. Imagine that you want to correct a set of errors in some files and push these corrections immediately. The problem

is that such a push will also include the latest, yet unfinished files that you have committed.

A better organization of your work would be to keep the latest, ongoing developments separate from the more official and stable version of the files. This is easily achieved by creating a separate branch where new developments takes place:

```
Terminal> git branch newstuff      # create and switch to new branch
Terminal> # extend and modify files...
Terminal> git commit -am 'Modified ... Added a file on ...'
Terminal> git checkout master      # swith back to master
Terminal> # correct errors
Terminal> git push origin master
Terminal> git checkout newstuff    # switch to other branch
Terminal> git merge master         # keep branch up-to-date w/master
Terminal> # continue development work...
Terminal> git commit -am 'More modifications of ...'
```

At some point, your developments in `newstuff` are mature enough to be incorporated in the master branch:

```
Terminal> git checkout newstuff
Terminal> git merge master         # synchronize newstuff w/master
Terminal> git checkout master
Terminal> git merge newstuff       # synchronize master w/newstuff
```

You no longer need the `newstuff` branch and can delete it:

```
Terminal> git branch -d newstuff
```

It is not possible to switch branches unless you have committed the files in the current branch. If your work on some files is in a mess and you want to change to another branch or fix other files in the current branch, a "global" commit affecting all files might be immature. Then the `git stash` command is handy. It records the state of your files and sets you back to the state of the last commit in the current branch. With `git stash apply` you get the files in this branch back to the state when you did the last `git stash`.

Let us explain a typical case. Suppose you have performed some extensive edits in some files and then you are suddenly interrupted. You need to fix some typos in some other files, commit the changes, and push. The problem is that many files are in an unfinished state - in hindsight you realize that those files should have been modified in a separate branch. It is not too late to create that branch! First run `git stash` to get the files back to the state they were at the last commit. Then run `git stash branch newstuff` to create a new branch `newstuff` containing the state of the files when you did the (last) `git stash` command. Stashing used this way is a convenient technique to move some immature edits after the last commit out in a new branch for further experimental work.

Warning.

You can get the stashed files back by `git stash apply`. It is possible to multiple `git stash` and `git stash apply` commands. However, it is easy to run into trouble with multiple stashes, especially if they occur in multiple

branches, as it becomes difficult to recognize which stashes that belong to which branch. A good advice is therefore to do `git stash only once` to get back to a clean state and then move the unfinished messy files to a separate branch with `git stash branch newstuff`.

5.4 Replacing pull by fetch and merge

The `git pull` command actually performs two steps that are sometimes advantageous to run separately. First, a `git fetch` is run to fetch new files from the repository, and thereafter a `git merge` command is run to merge the new files with your local version of the files. While `git pull` tries to do a lot and be smart in the merge, very often with success, the merge step may occasionally lead to trouble. That is why it is recommended to run a `git merge` separately, especially if you work with branches.

To fetch files from your repository at GitHub, which usually has the nickname `origin`, you write

```
Terminal> git fetch origin
```

You now have the possibility to check out in detail what the differences are between the new files and local ones:

```
Terminal> git diff origin/master
```

This command produces comparisons of the files in the current local branch and the `master` branch at `origin` (the GitHub repo). In this way you can exactly see the differences between branches. It also gives you an overview of what others have done with the files. When you are ready to merge in the new files from the `master` branch of `origin` with the files in the current local branch, you say

```
Terminal> git merge origin/master
```

Configuring the `git diff` command. The `git diff` command launches by default the Unix `diff` tool in the terminal window. Many users prefer to use other diff tools, and the desired one can be specified in your `~/.gitconfig` file. However, a much recommended approach is to wrap a shell script around the call to the diff program, because `git diff` actually calls the diff program with a series of command-line arguments that will confuse diff programs that take the names of the two files to be compared as arguments. In `~/.gitconfig` you specify a script to do the diff:

```
[diff]
external = ~/bin/git-diff-wrapper.sh
```

It remains to write the `git-diff-wrapper.sh` script. The 2nd and 5th command-line arguments passed to this script are the name of the files to be compared in the diff. A typical script may therefore look like

```
#!/bin/sh

diff "$2" "$5" | less
```

Here we use the standard (and quite primitive) Unix `diff` program, but we can replace `diff` by, e.g., `diffuse`, `kdiff3`, `xxdiff`, `meld`, `pdiff`, or others. With a Python script you can easily check for the extensions of the files and use different diff tools for different types of files, e.g., `latexdiff` for \LaTeX files and `pdiff` for pure text files.

Replacing all your files with those in the repo.

Occasionally it becomes desirable to replace *all* files in the local repo with those at the file hosting service. One possibility is removing your repo and cloning again, or use the Git commands

```
Terminal> git fetch --all
Terminal> git reset --hard origin/master
```

5.5 Team work with forking and pull requests

In small collaboration teams it is natural that everyone has push access to the repo. On GitHub this is known as the *Shared Repository Model*. As teams grow larger, there will usually be a few people in charge who should approve changes to the files. Ordinary team members will in this case not clone a repo and push changes, but instead *fork* the repo and send *pull requests*, which constitutes the *Fork and Pull Model*.

Say you want to fork the repo `https://github.com/somebody/proj1.git`. The first step is to press the *Fork* button on the project page for the `somebody/proj1` project on GitHub. This action creates a new repo `proj1`, known as the forked repo, on your GitHub account. Clone the fork as you clone any repo:

```
Terminal> git clone https://github.com/user/proj1.git
```

When you do `git push origin master`, you update your fork. However, the original repo is usually under development too, and you need to pull from that one to stay up to date. A `git pull origin master` pulls from `origin` which is your fork. To pull from the original repo, you create a name `upstream`,

```
Terminal> git remote add upstream \
    https://github.com/somebody/proj1.git
```

You can now do `git pull upstream` to get the most recent files in the original repo. However, it is not recommended to do a pull this way because heavy development of the `sombbody/proj1` project may lead to serious merge problems. It is much better to replace the pull by a separate *fetch* and *merge*. The typical workflow is

```
Terminal> git fetch upstream          # get new version of files
Terminal> git merge upstream/master   # merge with yours
Terminal> # Your files are up to date - ready for editing
Terminal> git commit -am 'Description...'
Terminal> git push origin master      # store changes in your fork
```

At some point you would like to push your changes back to the original repo `somebody/proj1`. This is done by a [pull request](#). Make sure you have

selected the right branch on the project page of your forked project. Press the *Pull Request* button and fill out the form that pops up. Trusted people in the `somebody/proj1` project will now review your changes and if they are approved, your files are merged into the original repo. If not, there are tools for keeping a dialog about how to proceed.

Also in small teams where everyone has push access, the fork and pull request model is beneficial for reviewing files before the repo is actually updated with new contributions.

5.6 Git workflows

Although the purpose of these notes is just to get the reader started with Git, it must be mentioned that there are advanced features of Git that have led to very powerful workflows with files and people, especially for software development. There is an official [Git workflow model](#) that outlines the basic principles, but it can be quite advanced for those with modest Git knowledge. A more detailed explanation of a recommended workflow for beginners is given in the [developer instructions](#) for the software package PETSc. This is highly suggested reading. The associated "quick summary" of Git commands for their workflow is also useful.

5.7 More documentation on Git

- [Web course on Git](#)
- [Everyday GIT With 20 Commands Or So](#)
- [Git top 10 tutorials](#)
- [Lars Vogel's Git Tutorial](#)
- [Git Community Book](#) (*explains* Git very well)
- [Git for Designers](#) (aimed a people with no previous knowledge of version control systems)
- [Git Magic: Basic Tricks](#)
- [The official Git Tutorial](#)
- [Git Tutorial Video](#) on YouTube
- [Git Questions](#)
- [Git Reference](#) (can also be used as a tutorial on Git)
- [Git User Manual](#)
- [Git home page](#)
- [Git/GitHub GUIs on Windows and Mac](#)

A Working with multiple GitHub accounts

Working against different GitHub accounts is easy if each project you work with on each account adds you as a collaborator. The term "you" here means your primary username on GitHub. My strong recommendation is to always check out a project using your primary GitHub username.

Occasionally you want to create a new GitHub account, say for a project XYZ. For such a non-personal account, do *not* provide an SSH key of any particular user. The reason is that this user will then get two GitHub identities, and switching between these identities will require some special tweakings. Just forget about the SSH key for a project account and add collaborators to repos using each collaborators personal GitHub username.

If you really need to operate the XYZ account as a personal account, you must provide an SSH key that is different from any other key at any other GitHub account (you will get an error message if you try to register an already registered SSH key, but it is possible to get around the error message by providing an `id_rsa.pub` key on one account and an `id_dsa.pub` on another - that will cause trouble). A recipe for how to operate multiple GitHub accounts using multiple identities is provided [here](http://net.tutsplus.com/tutorials/tools-and-tips/how-to-work-with-github-and-multiple-accounts/):

<http://net.tutsplus.com/tutorials/tools-and-tips/how-to-work-with-github-and-multiple-accounts/>

To debug which identity that is used when you pull and push to GitHub accounts, you can first run

```
Terminal> ssh -Tv git@github.com
```

to see your current identity and which SSH key that was used to identify you. Typing

```
Terminal> ssh-add -l
```

lists all your SSH keys. The shown strings can be compared with the string in the SSH key field of any GitHub account.