# Debugging in Python

**Hans Petter Langtangen**[1,2]

[1]Center for Biomedical Computing, Simula Research Laboratory
[2]Department of Informatics, University of Oslo

May 8, 2014

## Contents

Testing a program to find errors usually takes much more time than to write the code. This appendix is devoted to tools and good habits for effective debugging. Section 1 describes the Python debugger, a key tool for examining the internal workings of a code, while Section 2 explains how solve problems and write software to simplify the debugging process.

## 1 Using a debugger

A debugger is a program that can help you to find out what is going on in a computer program. You can stop the execution at any prescribed line number, print out variables, continue execution, stop again, execute statements one by one, and repeat such actions until you have tracked down abnormal behavior and found bugs.

Here we shall use the debugger to demonstrate the program flow of the code `Simpson.py`[1] (which can integrate functions of one variable with the famous Simpson's rule). This development of this code is explained in Section 4.2. You are strongly encouraged to carry out the steps below on your computer to get a glimpse of what a debugger can do.

---

[1]`http://tinyurl.com/pwyasaa/funcif/Simpson.py`

**Step 1.** Go to the folder `src/funcif`[2] where the program `Simpson.py` resides.

**Step 2.** If you use the Spyder Integrated Development Environment, choose *Debug* on the *Run* pull-down menu. If you run your programs in a plain terminal window, start IPython:

———————————————— Terminal ————————————————
```
Terminal> ipython
```
————————————————————————————————————————

Run the program `Simpson.py` with the debugger on (`-d`):

```
In [1]: run -d Simpson.py
```

We now enter the debugger and get a prompt

```
ipdb>
```

After this prompt we can issue various debugger commands. The most important ones will be described as we go along.

**Step 3.** Type `continue` or just `c` to go to the first line in the file. Now you can see a printout of where we are in the program:

```
1---> 1 def Simpson(f, a, b, n=500):
      2     """
      3     Return the approximation of the integral of f
```

Each program line is numbered and the arrow points to the next line to be executed. This is called the *current line*.

**Step 4.** You can set a *break point* where you want the program to stop so that you can examine variables and perhaps follow the execution closely. We start by setting a break point in the `application` function:

```
ipdb> break application
Breakpoint 2 at /home/.../src/funcif/Simpson.py:30
```

You can also say `break X`, where `X` is a line number in the file.

**Step 5.** Continue execution until the break point by writing `continue` or `c`. Now the program stops at line 31 in the `application` function:

```
ipdb> c
> /home/.../src/funcif/Simpson.py(31)application()
2     30 def application():
---> 31     from math import sin, pi
      32     print 'Integral of 1.5*sin^3 from 0 to pi:'
```

———————————————————————

[2]`http://tinyurl.com/pwyasaa/funcif`

**Step 6.** Typing `step` or just `s` executes one statement at a time:

```
ipdb> s
> /home/.../src/funcif/Simpson.py(32)application()
     31     from math import sin, pi
---> 32     print 'Integral of 1.5*sin^3 from 0 to pi:'
     33     for n in 2, 6, 12, 100, 500:

ipdb> s
Integral of 1.5*sin^3 from 0 to pi:
> /home/.../src/funcif/Simpson.py(33)application()
     32     print 'Integral of 1.5*sin^3 from 0 to pi:'
---> 33     for n in 2, 6, 12, 100, 500:
     34         approx = Simpson(h, 0, pi, n)
```

Typing another `s` reaches the call to `Simpson`, and a new `s` steps *into* the function `Simpson`:

```
ipdb> s
--Call--
> /home/.../src/funcif/Simpson.py(1)Simpson()
1---> 1 def Simpson(f, a, b, n=500):
     2     """
     3     Return the approximation of the integral of f
```

Type a few more `s` to step ahead of the `if` tests.

**Step 7.** Examining the contents of variables is easy with the `print` (or `p`) command:

```
ipdb> print f, a, b, n
<function h at 0x898ef44> 0 3.14159265359 2
```

We can also check the type of the objects:

```
ipdb> whatis f
Function h
ipdb> whatis a
<type 'int'>
ipdb> whatis b
<type 'float'>
ipdb> whatis n
<type 'int'>
```

**Step 8.** Set a new break point in the `application` function so that we can jump directly there without having to go manually through all the statements in the `Simpson` function. To see line numbers and corresponding statements around some line with number X, type `list X`. For example,

```
ipdb> list 32
     27 def h(x):
     28     return (3./2)*sin(x)**3
```

```
      29
      30 from math import sin, pi
      31
2     32 def application():
      33     print 'Integral of 1.5*sin^3 from 0 to pi:'
      34     for n in 2, 6, 12, 100, 500:
      35         approx = Simpson(h, 0, pi, n)
      36         print 'n=%3d, approx=%18.15f, error=%9.2E' % \
      37                 (n, approx, 2-approx)
```

We set a line break at line 35:

```
ipdb> break 35
Breakpoint 3 at /home/.../src/funcif/Simpson.py:35
```

Typing `c` continues execution up to the next break point, line 35.

**Step 9.** The command `next` or `n` is like `step` or `s` in that the current line is executed, but the execution does not step into functions, instead the function calls are just performed and the program stops at the next line:

```
ipdb> n
> /home/.../src/funcif/Simpson.py(36)application()
3     35         approx = Simpson(h, 0, pi, n)
--->  36         print 'n=%3d, approx=%18.15f, error=%9.2E' % \
      37                 (n, approx, 2-approx)
ipdb> print approx, n
1.9891717005835792 6
```

**Step 10.** The command `disable X Y Z` disables break points with numbers X, Y, and Z, and so on. To remove our three break points and continue execution until the program naturally stops, we write

```
ipdb> disable 1 2 3
ipdb> c
n=100, approx= 1.999999902476350, error= 9.75E-08
n=500, approx= 1.999999999844138, error= 1.56E-10

In [2]:
```

At this point, I hope you realize that a debugger is a very handy tool for monitoring the program flow, checking variables, and thereby understanding why errors occur.

# 2 How to debug

Most programmers will claim that writing code consumes a small portion of the time it takes to develop a program: the major portion of the work concerns testing the program and finding errors.

*Debugging is twice as hard as writing the code in the first place. Therefore, if you write the code as cleverly as possible, you are, by definition, not smart enough to debug it.* Brian W. Kernighan, computer scientist, 1942-.

Newcomers to programming often panic when their program runs for the first time and aborts with a seemingly cryptic error message. How do you approach the art of debugging? This appendix summarizes some important working habits in this respect. Some of the tips are useful for problem solving in general, not only when writing and testing Python programs.

## 2.1   A recipe for program writing and debugging

**1. Understand the problem.**   Make sure that you really understand the task the program is supposed to solve. We can make a general claim: if you do not understand the problem and the solution method, you will never be able to make a correct program. It may be argued that this claim is not entirely true: sometimes students with limited understanding of the problem are able to grab a similar program and guess at a few modifications, and actually get a program that works. But this technique is based on luck and not on understanding. The famous Norwegian computer scientist Kristen Nygaard (1926-2002) phrased it precisely: *Programming is understanding.* It may be necessary to read a problem description or exercise many times and study relevant background material before starting on the programming part of the problem solving process.

**2. Work out examples.**   Start with sketching one or more examples on input and output of the program. Such examples are important for controlling the understanding of the purpose of the program, and for verifying the implementation.

**3. Decide on a user interface.**   Find out how you want to get data into the program. You may want to grab data from the command-line, a file, or a dialog with questions and answers.

**4. Make algorithms.**   Identify the key tasks to be done in the program and sketch rough algorithms for these. Some programmers prefer to do this on a piece of paper, others prefer to start directly in Python and write Python-like code with comments to sketch the program (this is easily developed into real Python code later).

**5. Look up information.**   Few programmers can write the whole program without consulting manuals, books, and the Internet. You need to know and understand the basic constructs in a language and some fundamental problem solving techniques, but technical details can be looked up.

The more program examples you have studied (in this book, for instance), the easier it is to adapt ideas from an existing example to solve a new problem.

**6. Write the program.** Be extremely careful with what you write. In particular, compare all mathematical statements and algorithms with the original mathematical expressions.

In longer programs, do not wait until the program is complete before you start testing it, test parts while you write.

**7. Run the program.** If the program aborts with an error message from Python, these messages are fortunately quite precise and helpful. First, locate the line number where the error occurs and read the statement, then carefully read the error message. The most common errors (exceptions) are listed below.

`SyntaxError`: Illegal Python code.

```
File "somefile.py", line 5
  x = . 5
      ^
SyntaxError: invalid syntax
```

Often the error is precisely indicated, as above, but sometimes you have to search for the error on the previous line.

`NameError`: A name (variable, function, module) is not defined.

```
File "somefile.py", line 20, in <module>
  table(10)
File "somefile.py", line 16, in table
  value, next, error = L(x, n)
File "somefile.py", line 8, in L
  exact_error = log(1+x) - value_of_sum
NameError: global name 'value_of_sum' is not defined
```

Look at the last of the lines starting with `File` to see where in the program the error occurs. The most common reasons for a `NameError` are

- a misspelled name,

- a variable that is not initialized,

- a function that you have forgotten to define,

- a module that is not imported.


`TypeError`: An object of wrong type is used in an operation.

```
File "somefile.py", line 17, in table
  value, next, error = L(x, n)
File "somefile.py", line 7, in L
  first_neglected_term = (1.0/(n+1))*(x/(1.0+x))**(n+1)
TypeError: unsupported operand type(s) for +: 'float' and 'str'
```

Print out objects and their types (here: `print x, type(x), n, type(n)`), and you will most likely get a surprise. The reason for a `TypeError` is often far away from the line where the `TypeError` occurs.

`ValueError`: An object has an illegal value.

```
    File "somefile.py", line 8, in L
      y = sqrt(x)
  ValueError: math domain error
```

Print out the value of objects that can be involved in the error (here: `print x`).

`IndexError`: An index in a list, tuple, string, or array is too large.

```
    File "somefile.py", line 21
      n = sys.argv[i+1]
  IndexError: list index out of range
```

Print out the length of the list, and the index if it involves a variable (here: `print len(sys.argv), i`).

**8. Verify the results.** Assume now that we have a program that runs without error messages from Python. Before judging the results of the program, set precisely up a test case where you know the exact solution. This is in general quite difficult. In complicated mathematical problems it is an art to construct good test problems and procedures for providing evidence that the program works.

If your program produces wrong answers, start to *examine intermediate results*. Never forget that your own hand calculations that you use to test the program may be wrong!

**9. Use a debugger.** If you end up inserting a lot of `print` statements in the program for checking intermediate results, you might benefit from using a debugger as explained in Section 1.

Some may think that this list of nine points is very comprehensive. However, the recipe just contains the steps that you should always carry out when developing programs. Never forget that computer programming is a difficult task.

> *Program writing is substantially more demanding than book writing.*
> *Why is it so? I think the main reason is that a larger attention span*
> *is needed when working on a large computer program than when doing*
> *other intellectual tasks.* Donald Knuth [2, p. 18], computer scientist,
> 1938-.

## 2.2 Application of the recipe

Let us illustrate the points above in a specific programming problem: implementation of the Midpoint rule for numerical integration. The Midpoint rule for approximating an integral $\int_a^b f(x)dx$ reads

$$I = h \sum_{i=1}^{n} f(a + (i - \frac{1}{2})h), \quad h = \frac{b-a}{n}. \tag{1}$$

We just follow the individual steps in the recipe to develop the code.

**1. Understand the problem.** In this problem we must understand how to program the formula (1). Observe that we do not need to understand how the formula is derived, because we do not apply the derivation in the program. What is important, is to notice that the formula is an *approximation* of an integral. Comparing the result of the program with the exact value of the integral will in general show a discrepancy. Whether we have an approximation error or a programming error is always difficult to judge. We will meet this difficulty below.

**2. Work out examples.** As a test case we choose to integrate

$$f(x) = \sin^{-1}(x). \tag{2}$$

between 0 and $\pi$. From a table of integrals we find that this integral equals

$$\left[ x \sin^{-1}(x) + \sqrt{1-x^2} \right]_0^\pi. \tag{3}$$

The formula (1) gives an approximation to this integral, so the program will (most likely) print out a result different from (3). It would therefore be very helpful to construct a calculation where there are no approximation errors. Numerical integration rules usually integrate some polynomial of low order exactly. For the Midpoint rule it is obvious, if you understand the derivation of this rule, that a constant function will be integrated exactly. We therefore also introduce a test problem where we integrate $g(x) = 1$ from 0 to 10. The answer should be exactly 10.

*Input and output*: The input to the calculations is the function to integrate, the integration limits $a$ and $b$, and the $n$ parameter (number of intervals) in the formula (1). The output from the calculations is the approximation to the integral.

**3. Decide on a user interface.** We find it easiest at this beginning stage to program the two functions $f(x)$ and $g(x)$ directly in the program. We also specify the corresponding integration limits $a$ and $b$ in the program, but we read a common $n$ for both integrals from the command line. Note that this is not a flexible user interface, but it suffices as a start for creating a working program. A much better user interface is to read $f$, $a$, $b$, and $n$ from the command line, which will be done later in a more complete solution to the present problem.

**4. Make algorithms.** Like most mathematical programming problems, also this one has a generic part and an application part. The generic part is the formula (1), which is applicable to an arbitrary function $f(x)$. The implementation should reflect that we can specify any Python function `f(x)` and get it integrated. This principle calls for calculating (1) in a Python function where the input to the computation $(f, a, b, n)$ are arguments. The function heading can look as `integrate(f, a, b, n)`, and the value of (1) is returned.

The test part of the program consists of defining the test functions $f(x)$ and $g(x)$ and writing out the calculated approximations to the corresponding integrals.

A first rough sketch of the program can then be

```python
def integrate(f, a, b, n):
    # compute integral, store in I
    return I

def f(x):
...

def g(x):
...

# test/application part:
n = sys.argv[1]
I = integrate(g, 0, 10,  n)
print "Integral of g equals %g" % I
I = integrate(f, 0, pi,  n)
# calculate and print out the exact integral of f
```

The next step is to make a detailed implementation of the `integrate` function. Inside this function we need to compute the sum (1). In general, sums are computed by a `for` loop over the summation index, and inside the loop we calculate a term in the sum and add it to an accumulation variable. Here is the algorithm in Python code:

```python
s = 0
for i in range(1, n+1):
    s = s + f(a + (i-0.5)*h)
I = s*h
```

**5. Look up information.** Our test function $f(x) = \sin^{-1}(x)$ must be evaluated in the program. How can we do this? We know that many common mathematical functions are offered by the `math` module. It is therefore natural to check if this module has an inverse sine function. The best place to look for Python modules is the Python Standard Library[3] [1] documentation, which has a search facility. Typing *math* brings up a link to the `math` module, there we find `math.asin` as the function we need. Alternatively, one can use the command line utility `pydoc` and write `pydoc math` to look up all the functions in the module.

In this simple problem, we use very basic programming constructs and there is hardly any need for looking at similar examples to get started with the problem solving process. We need to know how to program a sum, though, via a `for` loop and an accumulation variable for the sum. Examples are found in Sections **??** and 1.8.

---

[3]`http://docs.python.org/2/library/`

**6. Write the program.** Here is our first attempt to write the program. You can find the whole code in the file `integrate_v1.py`[4].

```
def integrate(f, a, b, n):
    s = 0
    for i in range(1, n):
        s += f(a + i*h)
    return s

def f(x):
return asin(x)

def g(x):
return 1

# Test/application part
n = sys.argv[1]
I = integrate(g, 0, 10,  n)
print "Integral of g equals %g" % I
I = integrate(f, 0, pi,  n)
I_exact = pi*asin(pi) - sqrt(1 - pi**2) - 1
print "Integral of f equals %g (exact value is %g)' % \
  (I, I_exact)
```

**7. Run the program.** We try a first execution from IPython

```
In [1]: run integrate_v1.py
```

Unfortunately, the program aborts with an error:

```
    File "integrate_v1.py", line 8
      return asin(x)
           ^
  IndentationError: expected an indented block
```

We go to line 8 and look at that line and the surrounding code:

```
def f(x):
return asin(x)
```

Python expects that the return line is indented, because the function body must always be indented. By the way, we realize that there is a similar error in the `g(x)` function as well. We correct these errors:

```
def f(x):
    return asin(x)

def g(x):
    return 1
```

Running the program again makes Python respond with

---

[4]`http://tinyurl.com/pwyasaa/debug/integrate_v1.py`

```
    File "integrate_v1.py", line 24
      (I, I_exact)
                 ^
  SyntaxError: EOL while scanning single-quoted string
```

There is nothing wrong with line 24, but line 24 is a part of the statement
starting on line 23:

```
print "Integral of f equals %g (exact value is %g)' % \
      (I, I_exact)
```

A `SyntaxError` implies that we have written illegal Python code. Inspecting
line 23 reveals that the string to be printed starts with a double quote, but ends
with a single quote. We must be consistent and use the same enclosing quotes
in a string. Correcting the statement,

```
print "Integral of f equals %g (exact value is %g)" % \
      (I, I_exact)
```

and rerunning the program yields the output

```
  Traceback (most recent call last):
    File "integrate_v1.py", line 18, in <module>
      n = sys.argv[1]
  NameError: name 'sys' is not defined
```

Obviously, we need to import `sys` before using it. We add `import sys` and run
again:

```
  Traceback (most recent call last):
    File "integrate_v1.py", line 19, in <module>
      n = sys.argv[1]
  IndexError: list index out of range
```

This is a very common error: we index the list `sys.argv` out of range because
we have not provided enough command-line arguments. Let us use $n = 10$ in
the test and provide that number on the command line:

```
In [5]: run integrate_v1.py 10
```

We still have problems:

```
  Traceback (most recent call last):
    File "integrate_v1.py", line 20, in <module>
      I = integrate(g, 0, 10,  n)
    File "integrate_v1.py", line 7, in integrate
      for i in range(1, n):
  TypeError: range() integer end argument expected, got str.
```

It is the final `File` line that counts (the previous ones describe the nested
functions calls up to the point where the error occurred). The error message for
line 7 is very precise: the end argument to `range`, `n`, should be an integer, but it
is a string. We need to convert the string `sys.argv[1]` to `int` before sending it
to the `integrate` function:

11

```
n = int(sys.argv[1])
```

After a new edit-and-run cycle we have other error messages waiting:

```
Traceback (most recent call last):
  File "integrate_v1.py", line 20, in <module>
    I = integrate(g, 0, 10,  n)
  File "integrate_v1.py", line 8, in integrate
    s += f(a + i*h)
NameError: global name 'h' is not defined
```

The h variable is used without being assigned a value. From the formula (1) we see that $h = (b - a)/n$, so we insert this assignment at the top of the integrate function:

```
def integrate(f, a, b, n):
    h = (b-a)/n
    ...
```

A new run results in a new error:

```
Integral of g equals 9
Traceback (most recent call last):
  File "integrate_v1.py", line 23, in <module>
    I = integrate(f, 0, pi,  n)
NameError: name 'pi' is not defined
```

Looking carefully at all output, we see that the program managed to call the integrate function with g as input and write out the integral. However, in the call to integrate with f as argument, we get a NameError, saying that pi is undefined. When we wrote the program we took it for granted that pi was $\pi$, but we need to import pi from math to get this variable defined, before we call integrate:

```
from math import pi
I = integrate(f, 0, pi,  n)
```

The output of a new run is now

```
Integral of g equals 9
Traceback (most recent call last):
  File "integrate_v1.py", line 24, in <module>
    I = integrate(f, 0, pi,  n)
  File "integrate_v1.py", line 9, in integrate
    s += f(a + i*h)
  File "integrate_v1.py", line 13, in f
    return asin(x)
NameError: global name 'asin' is not defined
```

A similar error occurred: asin is not defined as a function, and we need to import it from math. We can either do a

12

```
from math import pi, asin
```

or just do the rough

```
from math import *
```

to avoid any further errors with undefined names from the `math` module (we will get one for the `sqrt` function later, so we simply use the last "import all" kind of statement).

There are still more errors:

```
Integral of g equals 9
Traceback (most recent call last):
  File "integrate_v1.py", line 24, in <module>
    I = integrate(f, 0, pi,  n)
  File "integrate_v1.py", line 9, in integrate
    s += f(a + i*h)
  File "integrate_v1.py", line 13, in f
    return asin(x)
ValueError: math domain error
```

Now the error concerns a wrong `x` value in the `f` function. Let us print out `x`:

```
def f(x):
    print x
    return asin(x)
```

The output becomes

```
Integral of g equals 9
0.314159265359
0.628318530718
0.942477796077
1.25663706144
Traceback (most recent call last):
  File "integrate_v1.py", line 25, in <module>
    I = integrate(f, 0, pi,  n)
  File "integrate_v1.py", line 9, in integrate
    s += f(a + i*h)
  File "integrate_v1.py", line 14, in f
    return asin(x)
ValueError: math domain error
```

We see that all the `asin(x)` computations are successful up to and including $x = 0.942477796077$, but for $x = 1.25663706144$ we get an error. A `math domain error` may point to a wrong $x$ value for $\sin^{-1}(x)$ (recall that the domain of a function specifies the legal $x$ values for that function).

To proceed, we need to think about the mathematics of our problem: Since $\sin(x)$ is always between $-1$ and $1$, the inverse sine function cannot take $x$ values outside the interval $[-1, 1]$. The problem is that we try to integrate $\sin^{-1}(x)$ from 0 to $\pi$, but only integration limits within $[-1, 1]$ make sense (unless we allow for complex-valued trigonometric functions). Our test problem is hence wrong from a mathematical point of view. We need to adjust the limits, say 0 to 1 instead of 0 to $\pi$. The corresponding program modification reads

13

```
I = integrate(f, 0, 1,  n)
```

We run again and get

```
Integral of g equals 9
0
0
0
0
0
0
0
0
0
Traceback (most recent call last):
  File "integrate_v1.py", line 26, in <module>
    I_exact = pi*asin(pi) - sqrt(1 - pi**2) - 1
ValueError: math domain error
```

It is easy to go directly to the `ValueError` now, but one should always examine
the output from top to bottom. If there is strange output before Python reports
an error, there may be an error indicated by our `print` statements. This is not
the case in the present example, but it is a good habit to start at the top of the
output anyway. We see that all our `print x` statements inside the `f` function
say that `x` is zero. This must be wrong - the idea of the integration rule is to
pick $n$ different points in the integration interval $[0, 1]$.

Our `f(x)` function is called from the `integrate` function. The argument to
`f`, `a + i*h`, is seemingly always 0. Why? We print out the argument and the
values of the variables that make up the argument:

```
def integrate(f, a, b, n):
    h = (b-a)/n
    s = 0
    for i in range(1, n):
        print a, i, h, a+i*h
        s += f(a + i*h)
    return s
```

Running the program shows that `h` is zero and therefore `a+i*h` is zero.

Why is `h` zero? We need a new `print` statement in the computation of `h`:

```
def integrate(f, a, b, n):
    h = (b-a)/n
    print b, a, n, h
    ...
```

The output shows that `a`, `b`, and `n` are correct. Now we have encountered a very
common error in Python version 2 and C-like programming languages: integer
division (see Section **??**). The formula $(1 - 0)/10 = 1/10$ is zero according to
integer division. The reason is that `a` and `b` are specified as `0` and `1` in the call
to `integrate`, and `0` and `1` imply `int` objects. Then `b-a` becomes an `int`, and
`n` is an `int`, causing an `int/int` division. We must ensure that `b-a` is `float` to
get the right mathematical division in the computation of `h`:

14

```
def integrate(f, a, b, n):
    h = float(b-a)/n
    ...
```

Thinking that the problem with wrong $x$ values in the inverse sine function is resolved, we may remove all the `print` statements in the program, and run again.

The output now reads

```
Integral of g equals 9
Traceback (most recent call last):
  File "integrate_v1.py", line 25, in <module>
    I_exact = pi*asin(pi) - sqrt(1 - pi**2) - 1
ValueError: math domain error
```

That is, we are back to the `ValueError` we have seen before. The reason is that `asin(pi)` does not make sense, and the argument to `sqrt` is negative. The error is simply that we forgot to adjust the upper integration limit in the computation of the exact result. This is another very common error. The correct line is

```
I_exact = 1*asin(1) - sqrt(1 - 1**2) - 1
```

We could have avoided the error by introducing variables for the integration limits, and a function for $\int f(x)dx$ would make the code cleaner:

```
a = 0; b = 1
def int_f_exact(x):
    return x*asin(x) - sqrt(1 - x**2)
I_exact = int_f_exact(b) - int_f_exact(a)
```

Although this is more work than what we initially aimed at, it usually saves time in the debugging phase to do things this proper way.

Eventually, the program seems to work! The output is just the result of our two `print` statements:

```
Integral of g equals 9
Integral of f equals 5.0073 (exact value is 0.570796)
```

**8. Verify the results.** Now it is time to check if the numerical results are correct. We start with the simple integral of 1 from 0 to 10: the answer should be 10, not 9. Recall that for this particular choice of integration function, there is no approximation error involved (but there could be a small round-off error). Hence, there must be a programming error.

To proceed, we need to calculate some intermediate mathematical results by hand and compare these with the corresponding statements in the program. We choose a very simple test problem with $n = 2$ and $h = (10 - 0)/2 = 5$. The formula (1) becomes

$$I = 5 \cdot (1 + 1) = 10 \,.$$

Running the program with $n = 2$ gives

```
Integral of g equals 1
```

We insert some `print` statements inside the `integrate` function:

```
def integrate(f, a, b, n):
    h = float(b-a)/n
    s = 0
    for i in range(1, n):
        print 'i=%d, a+i*h=%g' % (i, a+i*h)
        s += f(a + i*h)
    return s
```

Here is the output:

```
i=1, a+i*h=5
Integral of g equals 1
i=1, a+i*h=0.5
Integral of f equals 0.523599 (exact value is 0.570796)
```

There was only one pass in the `i` loop in `integrate`. According to the formula, there should be $n$ passes, i.e., two in this test case. The limits of `i` must be wrong. The limits are produced by the call `range(1,n)`. We recall that such a call results in integers going from 1 up to `n`, but *not* including `n`. We need to include `n` as value of `i`, so the right call to `range` is `range(1,n+1)`.

We make this correction and rerun the program. The output is now

```
i=1, a+i*h=5
i=2, a+i*h=10
Integral of g equals 2
i=1, a+i*h=0.5
i=2, a+i*h=1
Integral of f equals 2.0944 (exact value is 0.570796)
```

The integral of 1 is still not correct. We need more intermediate results!

In our quick hand calculation we knew that $g(x) = 1$ so all the $f(a + (i - \frac{1}{2})h)$ evaluations were rapidly replaced by ones. Let us now compute all the $x$ coordinates $a + (i - \frac{1}{2})h$ that are used in the formula:

$$i = 1 : \ a + (i - \frac{1}{2})h = 2.5, \quad i = 2 : \ a + (i - \frac{1}{2})h = 7.5 \, .$$

Looking at the output from the program, we see that the argument to `g` has a different value - and fortunately we realize that the formula we have coded is wrong. It should be `a+(i-0.5)*h`.

We correct this error and run the program:

```
i=1, a+(i-0.5)*h=2.5
i=2, a+(i-0.5)*h=7.5
Integral of g equals 2
...
```

Still the integral is wrong. At this point you may give up programming, but the more skills you pick up in debugging, the more fun it is to hunt for errors! Debugging is like reading an exciting criminal novel: the detective follows different ideas and tracks, but never gives up before the culprit is caught.

Now we read the code more carefully and compare expressions with those in the mathematical formula. We should, of course, have done this already when writing the program, but it is easy to get excited when writing code and hurry for the end. This ongoing story of debugging probably shows that reading the code carefully can save much debugging time. (Actually, being extremely careful with what you write, and comparing all formulas with the mathematics, may be the best way to get more spare time when taking a programming course!)

We clearly add up all the $f$ evaluations correctly, but then this sum must be multiplied by $h$, and we forgot that in the code. The `return` statement in `integrate` must therefore be modified to

```
    return s*h
```

Eventually, the output is

```
    Integral of g equals 10
    Integral of f equals 0.568484 (exact value is 0.570796)
```

and we have managed to integrate a constant function in our program! Even the second integral looks promising!

To judge the result of integrating the inverse sine function, we need to run several increasing $n$ values and see that the approximation gets better. For $n = 2, 10, 100, 1000$ we get $0.550371, 0.568484, 0.570714, 0.570794$, to be compared to the exact value $0.570796$. (This is not the mathematically exact value, because it involves computations of $\sin^{-1}(x)$, which is only approximately calculated by the `asin` function in the `math` module. However, the approximation error is very small ($\sim 10^{-16}$).) The decreasing error provides evidence for a correct program, but it is not a strong proof. We should try out more functions. In particular, linear functions are integrated exactly by the Midpoint rule. We can also measure the speed of the decrease of the error and check that the speed is consistent with the properties of the Midpoint rule, but this is a mathematically more advanced topic.

The very important lesson learned from these debugging sessions is that you should start with a simple test problem where all formulas can be computed by hand. If you start out with $n = 100$ and try to integrate the inverse sine function, you will have a much harder job with tracking down all the errors.

**9. Use a debugger.** Another lesson learned from these sessions is that we needed many `print` statements to see intermediate results. It is an open question if it would be more efficient to run a debugger and stop the code at relevant lines. In an edit-and-run cycle of the type we met here, we frequently need to examine many numerical results, correct something, and look at all the intermediate results again. Plain `print` statements are often better suited for this massive output than the pure manual operation of a debugger, unless one writes a program to automate the interaction with the debugger.

The correct code for the implementation of the Midpoint rule is found in `integrate_v2.py`[5]. Some readers might be frightened by all the energy it took

---

[5] `http://tinyurl.com/pwyasaa/debug/integrate_v2.py`

to debug this code, but this is just the nature of programming. The experience of developing programs that finally work is very awarding.

> *People only become computer programmers if they're obsessive about details, crave power over machines, and can bear to be told day after day exactly how stupid they are.* Gregory J. E. Rawlins [3], computer scientist.

**Refining the user interface.** We briefly mentioned that the chosen user interface, where the user can only specify $n$, is not particularly user friendly. We should allow $f$, $a$, $b$, and $n$ to be specified on the command line. Since $f$ is a function and the command line can only provide strings to the program, we may use the `StringFunction` object from `scitools.std` to convert a string expression for the function to be integrated to an ordinary Python function (see Section 3.3). The other parameters should be easy to retrieve from the command line if Section 2 is understood. As suggested in Section 7, we enclose the input statements in a `try-except` block, here with a specific exception type `IndexError` (because an index in `sys.argv` out of bounds is the only type of error we expect to handle):

```
try:
    f_formula = sys.argv[1]
    a = eval(sys.argv[2])
    b = eval(sys.argv[3])
    n = int(sys.argv[4])
except IndexError:
    print 'Usage: %s f-formula a b n' % sys.argv[0]
    sys.exit(1)
```

Note that the use of `eval` allows us to specify `a` and `b` as `pi` or `exp(5)` or another mathematical expression.

With the input above we can perform the general task of the program:

```
from scitools.std import StringFunction
f = StringFunction(f_formula)
I = integrate(f, a, b, n)
print I
```

**Writing a test function.** Instead of having these test statements as a main program we follow the good habits of Section 9 and make a module with

- the `integrate` function,

- a `test_integrate` function for testing the `integrate` function's ability to exactly integrate linear functions,

- a `main` function for reading data from the command line and calling `integrate` for the user's problem at hand.

Any module should also have a test block, as well as doc strings for the module itself and all functions.

The `test_integrate` function can perform a loop over some specified `n` values and check that the Midpoint rule integrates a linear function exactly. As always, we must be prepared for round-off errors, so "exactly" means errors less than (say) $10^{-14}$. The relevant code becomes

```python
def test_integrate():
    """Check that linear functions are integrated exactly."""

    def g(x):
        return p*x + q    # general linear function

    def int_g_exact(x):  # integral of g(x)
        return 0.5*p*x**2 + q*x

    a = -1.2; b = 2.8    # "arbitrary" integration limits
    p = -2;   q = 10
    success = True        # True if all tests below are passed
    for n in 1, 10, 100:
        I = integrate(g, a, b, n)
        I_exact = int_g_exact(b) - int_g_exact(a)
        error = abs(I_exact - I)
        if error > 1E-14:
            success = False
    assert success
```

We have followed the programming standard that will make this test function automatically work with the nose test framework:

1. the name of the function starts with `test_`,

2. the function has no arguments,

3. checks of whether a test is passed or not are done with `assert`.

The `assert success` statement raises an `AssertionError` exception if `success` is false, otherwise nothing happens. The nose testing framework searches for functions whose name start with `test_`, execute each function, and record if an `AssertionError` is raised. It is overkill to use nose for small programs, but in larger projects with many functions in many files, nose can run all tests with a short command and write back a notification that all tests passed.

The `main` function is simply a wrapping of the main program given above. The test block may call or `test_integrate` function or `main`, depending on whether the user will test the module or use it:

```python
if __name__ == '__main__':
    if sys.argv[1] == 'verify':
        verify()
    else:
        # Compute the integral specified on the command line
        main()
```

Here is a short demo computing $\int_0^{2\pi}(\cos(x) + \sin(x))dx$ with the aid of the `integrate.py`[6] file:

---
Terminal
---
```
integrate.py 'cos(x)+sin(x)' 0 2*pi 10
-3.48786849801e-16
```
---

## 2.3   Getting help from a code analyzer

The tools PyLint[7] and Flake8[8] can analyze your code and point out errors and undesired coding styles. Before point 7 in the lists above, *Run the program*, it can be wise to run PyLint or Flake8 to be informed about problems with the code.

Consider the first version of the `integrate` code, `integrate_v1.py`[9]. Running Flake8 gives

---
Terminal
---
```
Terminal> flake8 integrate_v1.py
integrate_v1.py:7:1: E302 expected 2 blank lines, found 1
integrate_v1.py:8:1: E112 expected an indented block
integrate_v1.py:8:7: E901 IndentationError: expected an indented block
integrate_v1.py:10:1: E302 expected 2 blank lines, found 1
integrate_v1.py:11:1: E112 expected an indented block
```
---

Flake8 checks if the program obeys the official Style Guide for Python Code[10] (known as *PEP8*). One of the rules in this guide is to have two blank lines before functions and classes (a habit that is often dropped in this book to reduce the length of code snippets), and our program breaks the rule before the `f` and `g` functions. More serious and useful is the `expected an indented block` at lines 8 and 11. This error is quickly found anyway by running the programming.

PyLint does not a complete job before the program is free of syntax errors. We must therefore apply it to the `integrate_v2.py`[11] code:

---
Terminal
---
```
Terminal> pylint integrate_v2.py
C: 20, 0: Exactly one space required after comma
I = integrate(f, 0, 1,  n)
                      ^ (bad-whitespace)
W: 19, 0: Redefining built-in 'pow' (redefined-builtin)
C:  1, 0: Missing module docstring (missing-docstring)
W:  1,14: Redefining name 'f' from outer scope (line 8)
W:  1,23: Redefining name 'n' from outer scope (line 16)
```

---
[6]http://tinyurl.com/pwyasaa/debug/integrate.py
[7]http://www.pylint.org/
[8]https://flake8.readthedocs.org/en/2.0/
[9]http://tinyurl.com/pwyasaa/debug/integrate_v1.py
[10]http://www.python.org/dev/peps/pep-0008/
[11]http://tinyurl.com/pwyasaa/debug/integrate_v2.py

```
C:  1, 0: Invalid argument name "f" (invalid-name)
C:  1, 0: Invalid argument name "a" (invalid-name)
```

There is much more output, but let us summarize what PyLint does not like about the code:

1. Extra whitespace (after comma in a call to `integrate`)

2. Missing doc string at the beginning of the file

3. Missing doc strings in the functions

4. Same name `f` used as local variable in `integrate` and global function name in the `f(x)` function

5. Too short variable names: `a`, `b`, `n`, etc.

6. "Star import" of the form `from math import *`

In short programs where the one-to-one mapping between mathematical notation and the variable names is very important to make the code self-explanatory, this author thinks that only points 1-3 qualify for attention. Nevertheless, for larger non-mathematical programs all the style violations pointed out are serious and lead to code that is easier to read, debug, maintain, and use.

Running Flak8 on `integrate_v2.py` leads to only three problems: missing two blank lines before functions (not reported by PyLint) and doing `from math import *`. Flake8 complains in general a lot less than PyLint, but both are very useful during program development to readability of the code and remove errors.

# References

[1] Python Software Foundation. The Python standard library. `http://docs.python.org/2/library/`.

[2] D. E. Knuth. Theory and practice. *EATCS Bull.*, 27:14–21, 1985.

[3] G. J. E. Rawlins. *Slaves of the Machine: The Quickening of Computer Technology.* MIT Press, 1998.

# Index